

OS THEORY ASSIGNMENT - 3

NAME: ARPIT MOHAPATRA

ROLL NO: 1905383

1. Create a system call called `getppid()` and create a command called "mytest" where you need to display the process-id along with parent process-id. (use the help of `getpid()`)

Solution:

1) First, the system call number is added to the `syscall.h` header file(line 23)

```
#define SYS_mknod 17
#define SYS_unlink 18
#define SYS_link 19
#define SYS_mkdir 20
#define SYS_close 21
#define SYS_getppid 22
#define SYS_sps 23
```

2) Next, the function prototype is added to the `defs.h` header file.(line 123)

```
void sched(void);
void setproc(struct proc*);
void sleep(void*, struct spinlock*);
void userinit(void);
int wait(void);
void wakeup(void*);
void yield(void);
int getpid(void);
int getppid(void);
int sps(void);
```

3) The same is added to the `user.h` header file.(line 26)

```
// system calls
int fork(void);
int exit(void) __attribute__((noreturn));
int wait(void);
int pipe(int*);
int write(int, const void*, int);
int read(int, void*, int);
int close(int);
int kill(int);
int exec(char*, char**);
int open(const char*, int);
int mknod(const char*, short, short);
int unlink(const char*);
int fstat(int fd, struct stat*);
int link(const char*, const char*);
int mkdir(const char*);
int chdir(const char*);
int dup(int);
int getpid(void);
char* sbrk(int);
int sleep(int);
int uptime(void);
int getppid(void);
int sps(void);
```

4) Next, the function is defined in `sysproc.c` (line 93)

```
int
sys_getppid(void)
{
    return myproc()->parent->pid;
}
```

5) The system call is added to `usys.S` (line 32)

```

SYSCALL(fork)
SYSCALL(exit)
SYSCALL(wait)
SYSCALL(pipe)
SYSCALL(read)
SYSCALL(write)
SYSCALL(close)
SYSCALL(kill)
SYSCALL(exec)
SYSCALL(open)
SYSCALL(mknod)
SYSCALL(unlink)
SYSCALL(fstat)
SYSCALL(link)
SYSCALL(mkdir)
SYSCALL(chdir)
SYSCALL(dup)
SYSCALL(getpid)
SYSCALL(sbrk)
SYSCALL(sleep)
SYSCALL(uptime)
SYSCALL(getppid)
SYSCALL_sps

```

6) The function is also added to `syscall.c` (line 106 & 131)

```

extern int sys_uptime(void);
extern int sys_getppid(void);
extern int sys_sps(void);

static int (*syscalls[])(void) = {
[SYS_fork]    sys_fork,
[SYS_exit]    sys_exit,
[SYS_wait]    sys_wait,
[SYS_pipe]    sys_pipe,
[SYS_read]    sys_read,
[SYS_kill]    sys_kill,
[SYS_exec]    sys_exec,
[SYS_fstat]   sys_fstat,
[SYS_chdir]   sys_chdir,
[SYS_dup]     sys_dup,
[SYS_getpid]  sys_getpid,
[SYS_sbrk]    sys_sbrk,
[SYS_sleep]   sys_sleep,
[SYS_uptime]  sys_uptime,
[SYS_open]    sys_open,
[SYS_write]   sys_write,
[SYS_mknod]   sys_mknod,
[SYS_unlink]  sys_unlink,
[SYS_link]    sys_link,
[SYS_mkdir]   sys_mkdir,
[SYS_close]   sys_close,
[SYS_getppid] sys_getppid,
[SYS_sps]     sys_sps,
};

```

6) A new `mytest.c` file is created and the following is added, to use the new system call.

```

File Edit View Search Terminal Help
#include "types.h"
#include "stat.h"
#include "user.h"
#include "fcntl.h"

int
main(int argc, char *argv[])
{
    printf(1, "Process id: %d\n", getpid());
    printf(1, "Parent process id: %d\n", getppid());
    exit();
}

```

8) Finally, the `Makefile` is modified to compile the `mytest.c` user program which uses our new system call. (line 183 & 254)

```
# Prevent deletion of
# that disk image chan
# details:
# http://www.gnu.org/s
.PRECIOUS: %.o

UPROGS=\
    _cat\
    _echo\
    _forktest\
    _grep\
    _init\
    _kill\
    _ln\
    _ls\
    _mkdir\
    _rm\
    _sh\
    _stressfs\
    _usertests\
    _wc\
    _mytest\
    _ps\
    _zombie\
```

```
$(QEMO) -nographic $(QEMO_OPTS) -S $(QEMO_DB)

# CUT HERE
# prepare dist for students
# after running make dist, probably want to
# rename it to rev0 or rev1 or so on and then
# check in that version.

EXTRA=\
    mkfs.c ulib.c user.h cat.c echo.c forktest.c grep.c kill.c\
    ln.c ls.c mkdir.c rm.c stressfs.c usertests.c wc.c mytest.c ps.c zombie.c\
    printf.c umalloc.c\
    README dot-bochsrc *.pl toc.* runoff runoff1 runoff.list\
    .gdbinit.tmpl gdbutil\

dist:
    rm -rf dist
```

9) Output:

```

ballocc: write bitmap block at sector 58
qemu-system-i386 -nographic -drive file=
xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nls
init: starting sh
$ ls
.          1 1 512
..         1 1 512
README    2 2 2286
cat        2 3 13676
echo       2 4 12684
forktest   2 5 8112
grep       2 6 15552
init       2 7 13268
kill       2 8 12740
ln         2 9 12636
ls         2 10 14820
mkdir     2 11 12820
rm        2 12 12796
sh        2 13 23284
stressfs   2 14 13468
usertests  2 15 56396
wc         2 16 14216
mytest     2 17 12628
ps         2 18 12416
zombie     2 19 12460
console    3 20 0
$ mytest
Process id: 4
Parent process id: 2
$

```

2. Create a ps command that will display the following. You need to prepare a system call called sps(system processes) that will provide the following information. PID, PPID, Process name, process state then you try to display the following Your roll no, PID, PPID, Process name, process state, process creation time, size of process memory

Solution:

1. First, the system call number is added to the `syscall.h` header file(line 24)

```

#define SYS_mknod 17
#define SYS_unlink 18
#define SYS_link 19
#define SYS_mkdir 20
#define SYS_close 21
#define SYS_getppid 22
#define SYS_sps 23
~
~
~

```

2. Next, the function prototype is added to the `defs.h` header file.(line 124)

```

void sched(void);
void setproc(struct proc*);
void sleep(void*, struct spinlock*);
void userinit(void);
int wait(void);
void wakeup(void*);
void yield(void);
int getppid(void);
int sps(void);
// switch 5

```

3. The same is added to the `user.h` header file.(line 27)

```

// system calls
int fork(void);
int exit(void) __attribute__((noreturn));
int wait(void);
int pipe(int*);
int write(int, const void*, int);
int read(int, void*, int);
int close(int);
int kill(int);
int exec(char*, char**);
int open(const char*, int);
int mknod(const char*, short, short);
int unlink(const char*);
int fstat(int fd, struct stat*);
int link(const char*, const char*);
int mkdir(const char*);
int chdir(const char*);
int dup(int);
int getpid(void);
char* sbrk(int);
int sleep(int);
int uptime(void);
int getppid(void);
int sps(void);

```

4. Next, the function is defined in `sysproc.c` (line 99)

```

    return myproc()->parent->pid;
}

int
sys_sps(void)
{
    return sps();
}
"sysproc.c" 103L, 1196C

```

5. Start time property is added to `proc` struct in `proc.h` (line 42)

```

// Per-process state
struct proc {
    uint sz;                // Size of process memory (bytes)
    pde_t* pgdir;          // Page table
    char* kstack;           // Bottom of kernel stack for this process
    uint stime;             // Start time
    enum procstate state;   // Process state
    int pid;               // Process ID
    struct proc* parent;    // Parent process
    struct trapframe* tf;   // Trap frame for current syscall
    struct context* context; // switch() here to run process
    void* chan;            // If non-zero, sleeping on chan
    int killed;            // If non-zero, have been killed
}
"proc.h" 59L, 2310C

```

6. `sps()` function is added in `proc.c` with other changes to account for start time. (line 91 & 551)

```

86     return 0;
87
88 found:
89     p->state = EMBRYO;
90     p->pid = nextpid++;
91     p->stime = ticks;
92
93     release(&ptable.lock);
94

```

```

551 int sps(void)
552 {
553     struct proc *p;
554     char *states[] = {
555         [UNUSED] "Unused",
556         [EMBRYO] "Embryo",
557         [SLEEPING] "Sleeping",
558         [RUNNABLE] "Runnable",
559         [RUNNING] "Running",
560         [ZOMBIE] "Zombie",
561     };
562     acquire(&ptable.lock);
563     cprintf("ROLL\tPID\tPPID\tNAME\tSIZE\tSTART\tSTATE\n");
564     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
565         if(p->state != UNUSED) {
566             cprintf("1905383");
567             cprintf("\t%d", p->pid);
568             if(p->parent->pid/10 == 0) {
569                 cprintf("\t%d", p->parent->pid);
570             }
571             else {
572                 cprintf("\tNone");
573             }
574         }
575         cprintf("\t%s", p->name);
576         cprintf("\t%d", p->sz);
577         cprintf("\t%d", p->stime);
578         cprintf("\t%s", states[p->state]);
579         cprintf("\n");
580     }
581     release(&ptable.lock);
582     return 23;
583 }

```

7. The system call is added to `usys.S` (line 33)

```

File Edit View Search Terminal Help
#include "syscall.h"
#include "traps.h"

#define SYSCALL(name) \
.globl name; \
name: \
    movl $SYS_ ## name, %eax; \
    int $T_SYSCALL; \
    ret

SYSCALL(fork)
SYSCALL(exit)
SYSCALL(wait)
SYSCALL(pipe)
SYSCALL(read)
SYSCALL(write)
SYSCALL(close)
SYSCALL(kill)
SYSCALL(exec)
SYSCALL(open)
SYSCALL(mknod)
SYSCALL(unlink)
SYSCALL(fstat)
SYSCALL(link)
SYSCALL(mkdir)
SYSCALL(chdir)
SYSCALL(dup)
SYSCALL(getpid)
SYSCALL(sbrk)
SYSCALL(sleep)
SYSCALL(uptime)
SYSCALL(getppid)
SYSCALL(sps)

```

8. The function is also added to `syscall.c` (line 107 & 132)

```
static int (*syscalls[])(void) = {
[SYS_fork]      sys_fork,
[SYS_exit]      sys_exit,
[SYS_wait]      sys_wait,
[SYS_pipe]      sys_pipe,
[SYS_read]      sys_read,
[SYS_kill]      sys_kill,
[SYS_exec]      sys_exec,
[SYS_fstat]     sys_fstat,
[SYS_chdir]     sys_chdir,
[SYS_dup]       sys_dup,
[SYS_getpid]    sys_getpid,
[SYS_sbrk]      sys_sbrk,
[SYS_sleep]     sys_sleep,
[SYS_uptime]    sys_uptime,
[SYS_open]      sys_open,
[SYS_write]     sys_write,
[SYS_mknod]     sys_mknod,
[SYS_unlink]    sys_unlink,
[SYS_link]      sys_link,
[SYS_mkdir]     sys_mkdir,
[SYS_close]     sys_close,
[SYS_getppid]   sys_getppid,
[SYS_sps]       sys_sps,
};

void
syscall(void)
```

9. A new `ps.c` file is created and the following is added, to use the new system call.

```
File Edit View Search Terminal
#include "types.h"
#include "stat.h"
#include "user.h"

int main(void)
{
    sps();
    exit();
}
```

10. Finally, the `Makefile` is modified to compile the `ps.c` user program which uses our new system call. (line 184 & 254)

```
# Prevent deletion of
# that disk image chan
# details:
# http://www.gnu.org/s
.PRECIOUS: %.o

UPROGS=\
    _cat\
    _echo\
    _forktest\
    _grep\
    _init\
    _kill\
    _ln\
    _ls\
    _mkdir\
    _rm\
    _sh\
    _stressfs\
    _usertests\
    _wc\
    _mytest\
    _ps\
    _zombie\
```

```

$(QEMU) -nographic $(QEMUPYS) -S $(QEMUGDB)

# CUT HERE
# prepare dist for students
# after running make dist, probably want to
# rename it to rev0 or rev1 or so on and then
# check in that version.

EXTRA=\
mkfs.c ulib.c user.h cat.c echo.c forktest.c grep.c kill.c\
ln.c ls.c mkdir.c rm.c stressfs.c usertests.c wc.c mytest.c ps.c zombie.c\
printf.c umalloc.c\
README dot-bochsrc *.pl toc.* runoff runoff1 runoff.list\
.gdbinit.tmpl gdbutil\

dist:
rm -rf dist

```

11. Output:

```

350+1 records out
179620 bytes (180 kB, 175 KiB) copied, 0.00142921 s, 126 MB/s
qemu-system-i386 -nographic -drive file=fs.img,index=1,media=disk,format=raw...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32
init: starting sh
$ ls
.          1 1 512
..         1 1 512
README    2 2 2286
cat       2 3 13676
echo      2 4 12684
forktest  2 5 8112
grep      2 6 15552
init      2 7 13268
kill      2 8 12740
ln        2 9 12636
ls        2 10 14820
mkdir     2 11 12820
rm        2 12 12796
sh        2 13 23284
stressfs  2 14 13468
usertests 2 15 56396
wc        2 16 14216
mytest    2 17 12628
ps        2 18 12416
zombie    2 19 12460
console   3 20 0
$ ps
ROLL  PID  PPID  NAME  SIZE  START  STATE
1905383 1      None  init  12288  0      Sleeping
1905383 2        1    sh   16384  31     Sleeping
1905383 4        2    ps   12288  383    Running
          0        0  Unused

```